

# GENERATIVE ADVERSARIAL NETWORKS

# What are GANs?

## Generative Adversarial Networks

Generative Models

We try to learn the underlying the distribution  
from which our dataset comes from.

# What are GANs?

## Generative Adversarial Networks

Adversarial Training

GANs are made up of two competing networks (adversaries) that are trying to beat each other.

What are GANs?

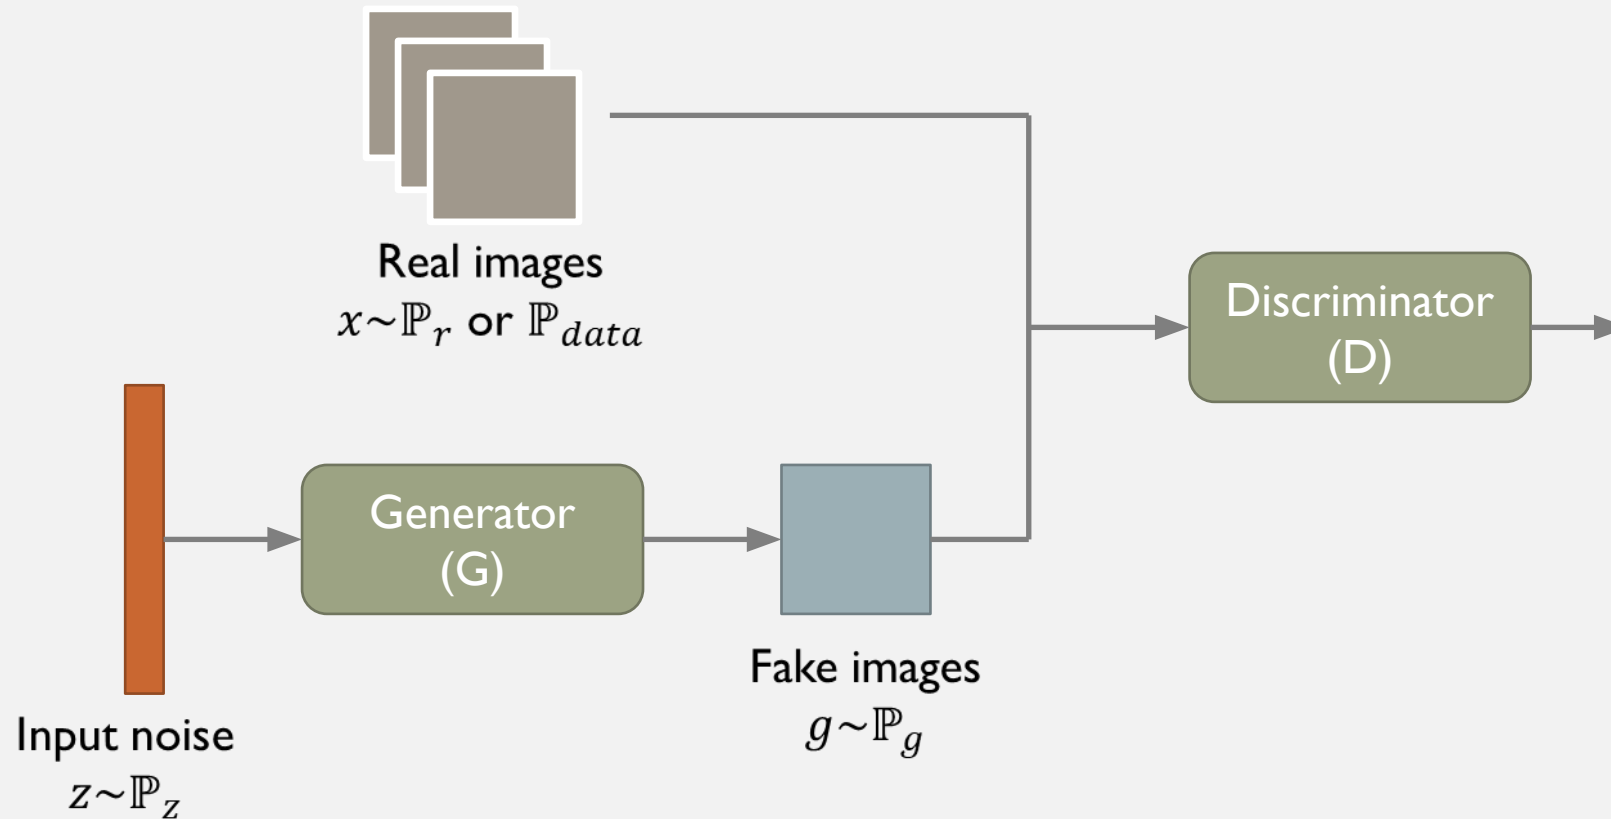
Generative Adversarial Networks  
Neural Networks

# GENERATIVE VS DISCRIMINATIVE

Goal of ML algorithms: “predict the probability of a sample belonging to class  $y$ , based on a set of features  $x$ ”. In other words:  $p(y | x)$ .

- **Discriminative** algorithms try to solve this by correlating features to labels.  
E.g. logistic regression, SVM
- **Generative** algorithms try to learn the distribution of each class.  
“How did we get  $x$ ?”  $\rightarrow p(x | y)$   
E.g. Naïve Bayes

# GENERATIVE ADVERSARIAL NETWORKS



# GENERATOR VS DISCRIMINATOR

- Generator's goal: produce “fake” data that will trick the discriminator
- Discriminator's goal: distinguish between real and fake data
- The two compete in a zero-sum game, which will hopefully improve the performance of both models.
- Very hard to train. Especially during early stages.

## IN MORE DETAILS...

• Original conception (Goodfellow et al.):

- $D$  and  $G$  play a two-player minimax game with value function  $V(D, G)$  :

$$\min_G \max_D V(D, G) = \mathbb{E}_{x \sim p_{data}(x)} [\log D(x)] + \mathbb{E}_{z \sim p_z(z)} [1 - \log D(G(z))]$$

$p_{data}$ : real data  
distribution

$D(x)$ : probability that  
 $x$  came from the data  
rather than  $p_g$

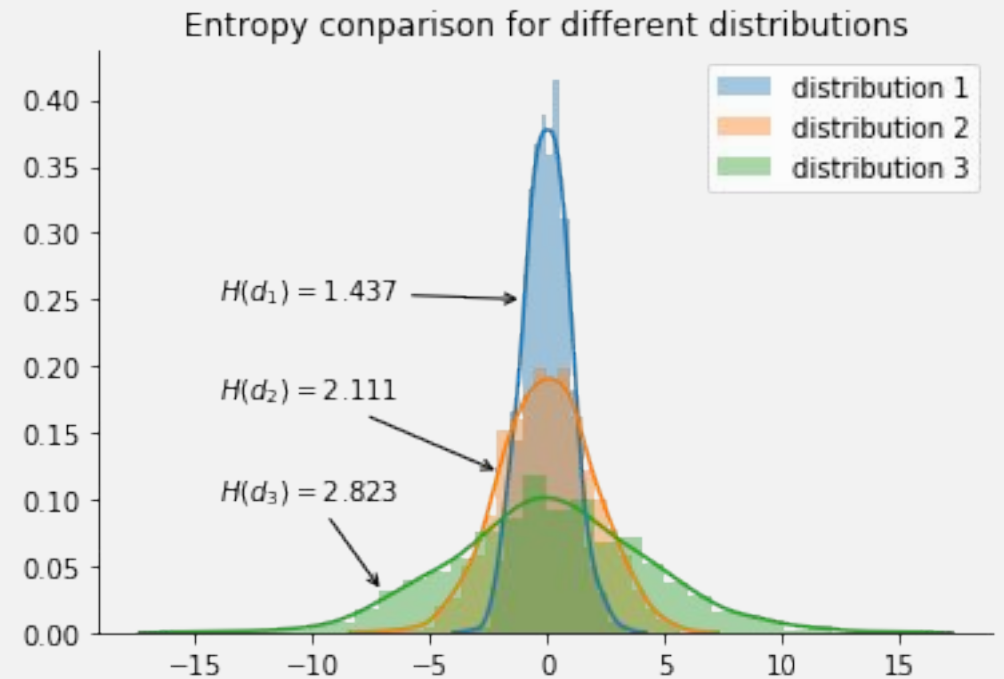
$p_z$ : input noise  
prior

- Discriminator's loss is essentially a **cross-entropy** loss.



# SHANNON ENTROPY


- $$H(P) = \sum_x P(x) \cdot \log\left(\frac{1}{P(x)}\right)$$
$$H(P) = -\mathbb{E}_{x \sim p}[\log P(x)]$$
- Bits needed to encode “message”  $x$  with probability  $P(x)$ .
- Entropy  $H$ : sum over all messages.
- Measure of “**impurity**” in data.
- **Low entropy** → data more uniform.



# CROSS-ENTROPY

- $$H(y, \hat{y}) = \sum_i y_i \cdot \log\left(\frac{1}{\hat{y}_i}\right)$$
$$H(y, \hat{y}) = -\mathbb{E}_y[\log \hat{y}]$$
- Bits needed to encode “message”  $i$  from distribution  $y$  with “symbols” from distribution  $\hat{y}$ .
- Measure “**relatedness**” between  $y$  and  $\hat{y}$ .
- **High cross-entropy**  $\rightarrow (y, \hat{y})$  **unrelated**.

# CROSS-ENTROPY INSIGHTS



|  | 1.437 | 1.748 | 2.184 |
|--|-------|-------|-------|
|  | 2.865 | 2.087 | 2.311 |
|  | 5.583 | 3.495 | 2.577 |

$$H(d_1, d_1) = H(d_1) = 1.437$$

$$H(d_2, d_2) = H(d_2) = 2.087$$

$$H(d_3, d_3) = H(d_3) = 2.577$$

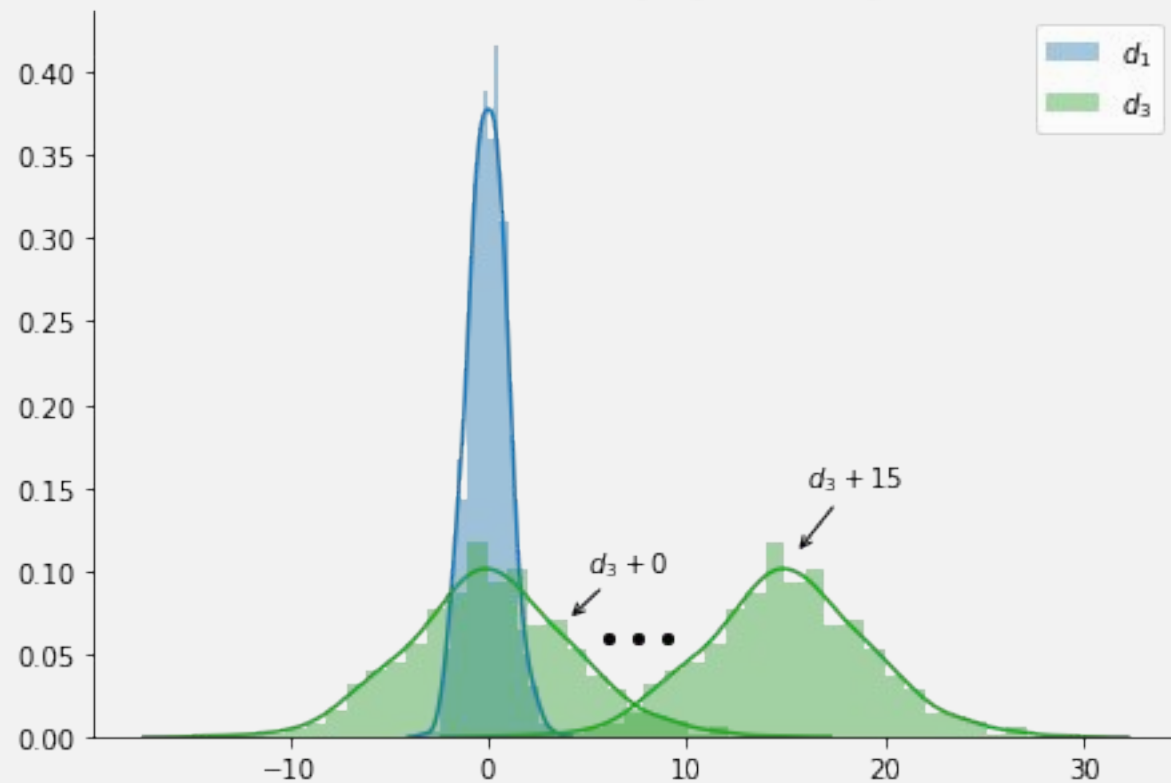
$$H(d_1, d_2) \neq H(d_2, d_1)$$

$$H(d_1, d_3) \neq H(d_3, d_1)$$


...

# CROSS-ENTROPY INSIGHTS

Distributions:  $d_1$ ,  $d_3$ , shifted  $d_3$



Cross-entropy between the two distributions



# KULLBACK-LEIBLER (KL) DIVERGENCE

- Measures **difference** between two probability distributions  $(y, \hat{y})$ .
- Essentially the **difference from entropy to cross-entropy**.

$$KL(y \parallel \hat{y}) = H(y, \hat{y}) - H(y)$$
$$KL(y \parallel \hat{y}) = \sum_i y_i \cdot \log \left( \frac{y_i}{\hat{y}_i} \right)$$

- How many more bits are required to encode messages from distribution  $y$  if we use symbols from  $\hat{y}$ .
- **Non symmetric**:  $KL(y \parallel \hat{y}) \neq KL(\hat{y} \parallel y)$ . Not a distance metric.
- Always takes **positive** values:  $H(y, \hat{y}) \geq H(y)$ .
- The **higher** the value, the more two distributions **differ**.

## JENSEN-SHANNON (JS) DIVERGENCE

- A derivative of KL divergence, that **can be used as a distance metric**.

$$JS(y || \hat{y}) = \frac{1}{2} KL(y || m) + \frac{1}{2} KL(\hat{y} || m)$$

where  $m = \frac{1}{2}(y + \hat{y})$

- **Symmetric.**
- **Positive.**
- The **higher** the value the **larger** the distance between  $y, \hat{y}$ .

# GAN LOSSES

- $\min_G \max_D V(D, G) = \mathbb{E}_{x \sim p_{data}(x)} [\log D(x)] + \mathbb{E}_{z \sim p_z(z)} [1 - \log D(G(z))]$

- Discriminator loss (sum two losses):

- Loss on “real” images:

$$- \mathbb{E}_{x \sim p_{data}(x)} [\log D(x)]$$

- Loss on “fake” images:

$$- \mathbb{E}_{z \sim p_z(z)} [1 - \log D(G(z))]$$

- Generator loss (**only relevant for “fake” images**):

$$\mathbb{E}_{z \sim p_z(z)} [1 - \log D(G(z))]$$

# TRAINING

- Initially:
  - Generator does not know how to produce realistic images.
  - Discriminator does not know how to separate the two.
- As training progresses:
  - G: starts to learn the distribution of real images.
  - D: becomes better at distinguishing real from fake.
- After successful training:
  - G: has learned to produce realistic images.
  - D: can't distinguish from the two .



## TRAINING DETAILS

- When you train the discriminator, **hold the generator values constant**; and when you train the generator, **hold the discriminator constant**. Each should train against a static adversary.
- Each side of the GAN can overpower the other.
  - If the discriminator is too good, it will return values so close to 0 or 1 that the generator will struggle to read the gradient.
  - If the generator is too good, it will persistently exploit weaknesses in the discriminator that lead to false negatives.

# GAN TRAINING ISSUES

- $$\min_G \max_D V(D, G) = \mathbb{E}_{x \sim p_{data}(x)} [\log D(x)] + \mathbb{E}_{z \sim p_z(z)} [1 - \log D(G(z))]$$
- Hard to achieve Nash Equilibrium (each player plays independently; cannot ensure convergence).
- **Vanishing Gradient:**  
As D becomes better (makes more confident predictions), G loss approaches 0  $\rightarrow$  no gradient.  
Perfect discriminator exists (proven).
- If D behaves badly, G cannot gather accurate feedback.
- JS divergence fails to provide a meaningful metric when the two distributions are disjointed.
- Highly sensitive to hyperparameter selection.
- **Mode collapse:**  
Multimodal distribution “modes” collapse into a single “mode”.

# MODE COLLAPSE

- Consider the extreme case that G is trained without updating D. The generated images will converge towards the optimal image  $x^*$ , which maximizes D's uncertainty:  $x^* = \operatorname{argmax}_x D(x)$ .
- The mode collapses to a single point, the gradient associated with  $z$  becomes 0.



- D will then push to exploit the next best mode. G will soon follow and **the two will be stuck in an overfit** state where both will try to exploit the other's short term weaknesses.
- D is more likely to overfit than G.

# SOME WORKAROUNDS TO GAN'S TRAINING ISSUES...

- **Minibatch discrimination:**

- Separate real/fake into different batches.
- Compute similarity within batch.
- Feed this to the discriminator.
- Mode drops -> similarity increases.
- D can spot fake images from this parameter and penalize G.

- **One-sided label smoothing:**

- Penalize D when prediction for real is high. E.g.  $D(x) > 0.9$ .
- Avoid overconfidence.

- **Add noise to stabilize the model.**

- **Alternative generator cost function:**

Replace  $\mathbb{E}_{z \sim p_z(z)}[1 - \log D(G(z))]$  with  $-\mathbb{E}_{z \sim p_z(z)}[\log D(G(z))]$

- More resistant to vanishing gradients

## WASSERSTEIN GAN (I/4)

- Most of GAN's training issues emanate from the **cost function**.
- Replace JS with “Wasserstein distance”.
- **Distance metric between distributions.**
- How much “earth” we need to move to reach from one distribution to the other.
- Continuous everywhere and differentiable almost everywhere (not true for KL, JS).

$$W(\mathbb{P}_r, \mathbb{P}_g) = \inf_{\gamma \in \Pi(\mathbb{P}_r, \mathbb{P}_g)} \mathbb{E}_{(x,y) \sim \gamma} [ ||x - y|| ]$$

where  $\Pi(\mathbb{P}_r, \mathbb{P}_g)$  is the set of all joint distributions  $(x, y)$  whose marginals are  $\mathbb{P}_r, \mathbb{P}_g$ .

## WASSERSTEIN GAN (2/4)

- The previous cost function depends on the optimal “transport plan”  $\gamma$ , which is tricky to compute.
- After some assumptions and by using the Kantorovich-Rubinstein duality, we can simplify the calculation to:

$$W(\mathbb{P}_r, \mathbb{P}_\theta) = \sup_{\|f\|_L \leq 1} \mathbb{E}_{x \sim \mathbb{P}_r}[f(x)] - \mathbb{E}_{x \sim \mathbb{P}_\theta}[f(x)]$$

where  $f$  is a 1-Lipschitz function.

- In practice we build a deep neural network to **approximate** this function 😊. This new network is very similar to the Discriminator, which we will call the “**critic**” to reflect its new role.

## WASSERSTEIN GAN (3/4)

- New gradients:

- **Critic:**

$$\nabla_w \frac{1}{m} \sum_{i=1}^m [f(x^{(i)}) - f(G(z^{(i)}))]$$

- **Generator:**

$$\nabla_{\theta} \frac{1}{m} \sum_{i=1}^m [-f(G(z^{(i)}))]$$

- Unfortunately  $f$  needs to be a 1-Lipschitz function. To enforce this constraint the authors simply **clipped**  $f$  if it exceeded an arbitrary range.

# WASSERSTEIN GAN (4/4)

## Pros:

- Much more **meaningful loss** (correlation between loss value and image quality).
- Increased training **stability**:
  - Less likely to collapse.
  - **Generator can still learn when critic performs well.**
  - Allows training to **optimality**.

## Cons (all have to do with weight clipping):

- Clipping parameter is very sensitive.
- Decreases model capacity.
- Slows down training.



## GRADIENT PENALTY (WGAN-GP)

- Done instead of weight clipping.
- A 1-Lipshitz function has a gradient norm of 1 almost everywhere under  $\mathbb{P}_r, \mathbb{P}_\theta$ .
- The model is **penalized** if the gradient norm strays away from 1.


$$L = \underbrace{\mathbb{E}_{\tilde{x} \sim \mathbb{P}_g}[D(\tilde{x})] - \mathbb{E}_{x \sim \mathbb{P}_r}[D(x)]}_{\text{critic loss}} + \underbrace{\lambda \mathbb{E}_{\hat{x} \sim \mathbb{P}_{\tilde{x}}}[(\|\nabla_{\hat{x}} D(\hat{x})\|_2 - 1)^2]}_{\text{gradient penalty}}$$

where  $\hat{x}$  is sampled from  $\tilde{x}$  and  $x$ .

- **Avoid** Batch Normalization (creates correlation between samples within batch, impacts the effectiveness of GP).
- Makes more **stable** training and will **converge** better and faster.
- Allows for the use of more **complex** models as G, D (e.g. ResNet architecture).

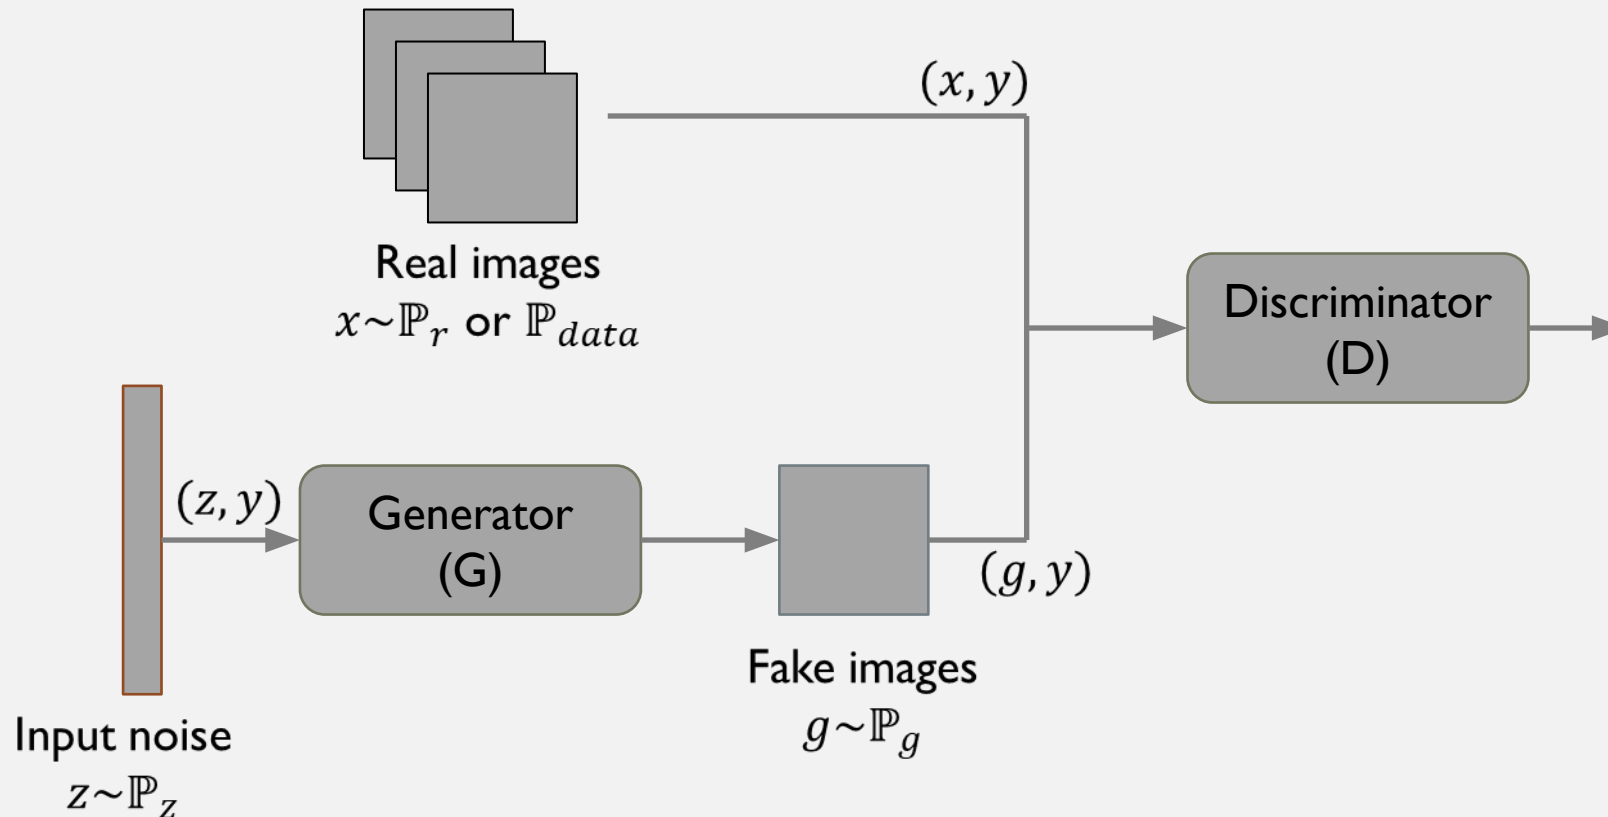
# DEEP CONVOLUTIONAL GAN (DCGAN)

- Original GANs were MLP models.
- Can build GANs with **convolutional layers**.
- Guidelines:
  - Replace pooling layers (which destroy spatial information) with convolutional stride.
  - Use transposed convolutions for up-sampling.
  - Eliminate FC layers.
  - Use Batch Normalization.



# CONDITIONAL GAN (CGAN)

- Both generator and discriminator also accept the image labels as their input.



## LEAST SQUARES GAN (LSGAN)

- Defines a new cost function which help get a smoother gradient everywhere.

$$\begin{aligned}\min_D V_{\text{LSGAN}}(D) &= \frac{1}{2} \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}(\mathbf{x})} [(D(\mathbf{x}) - b)^2] + \frac{1}{2} \mathbb{E}_{\mathbf{z} \sim p_{\mathbf{z}}(\mathbf{z})} [(D(G(\mathbf{z})) - a)^2] \\ \min_G V_{\text{LSGAN}}(G) &= \frac{1}{2} \mathbb{E}_{\mathbf{z} \sim p_{\mathbf{z}}(\mathbf{z})} [(D(G(\mathbf{z})) - c)^2],\end{aligned}$$

# ENERGY BASED GAN (EBGAN) AND BOUNDARY EQUILIBRIUM GAN (BEGAN)

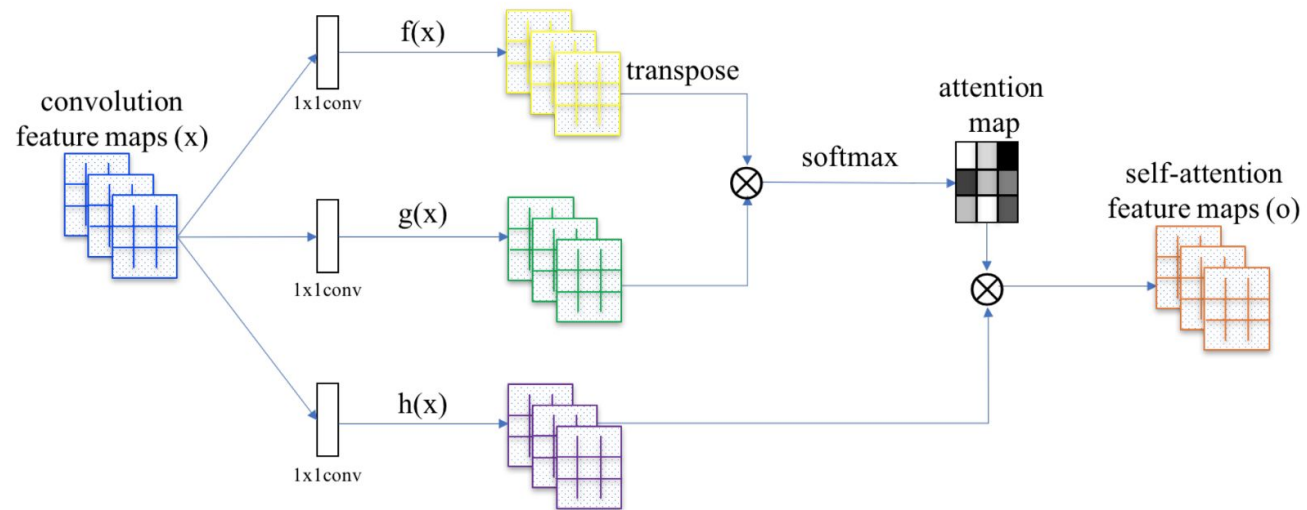
- Both replace D with an **AutoEncoder**.
- **EBGAN:**
  - Measure reconstruction error (MSE)
  - Motivates GAN to have broader goals and avoid greedy optimization.
  - Ensures GAN generates images with features found in natural images.
- **BEGAN:**
  - Wasserstein distance
  - Loss function has two goals (hyperparameter  $\gamma$  to balance them):
    - Good critic (helps diversity).
    - Good reconstructor (helps image quality).

# SPECTRAL NORMALIZATION

- Popular discriminator regularizer.
- Can enforce Lipschitz constraint.

## SELF-ATTENTIONAL GAN (SEGAN)

- **State-of-the-art!**
- Refines each spatial location with an extra term computed by the self-attentional mechanism.
- Can be used on both G and D.
- Spectral Normalization is used to stabilize the GAN.



## EVALUATION METRICS (1/2)

### Inception Score:

- Entropy can be viewed as randomness.
- When training a classifier we want  $p(y | x)$  to be **highly predictable**.
- Use an **Inception** network to classify the images and predict  $p(y | x)$ . This reflects on the quality of the images.
- If the generated images are diverse the data distribution of  $y$  should be **uniform**:

$$p(y) = \int_z p(y | x = G(z)) dz$$

- Inception score is computed as:

$$IS(G) = \exp \left( \mathbb{E}_{x \sim p_g} [KL(\overbrace{p(y | x)}^{\text{High } p(y | x) \text{ requires high quality images}} || \underbrace{p(y)}_{\text{Large } p(y) \text{ requires high diversity between classes}})] \right)$$

Large  $p(y)$  requires high diversity between classes



## EVALUATION METRICS (2/2)

### Fréchet Inception Distance (FID)

- Use **Inception** Network to **extract features from an intermediate layer**.
- Model these with a multivariate Gaussian distribution with mean  $\mu$  and covariance  $\Sigma$ .
- FID between real images  $x$  and generated  $g$  images is computed as:

$$FID(x, g) = || \mu_x - \mu_g ||_2^2 + Tr \left( \Sigma_x + \Sigma_g - 2(\Sigma_x \Sigma_g)^{\frac{1}{2}} \right)$$

- **Lower** FID  $\rightarrow$  better image quality and diversity.
- Sensitive to mode collapse (increases with missing modes).
- More robust to noise than IS.
- Better measure for image diversity.