



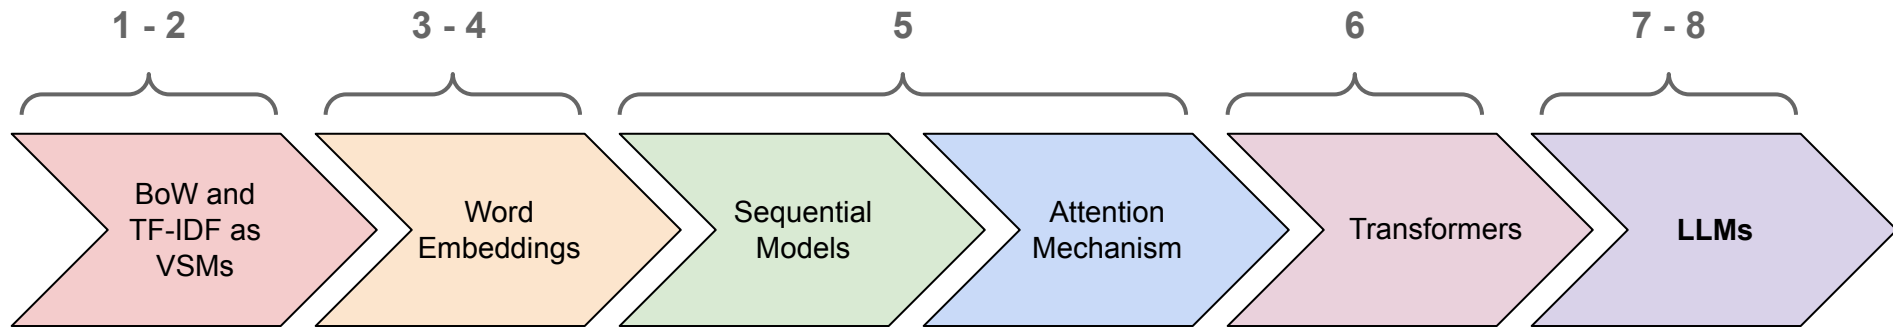
# Code.Hub

The first Hub for Developers

## Large Language Models

Thanos Tagaris

# NLP timeline up till today...



- The first direction the research community took was to scale their Transformers as much as they could.
- These large-scale models became really proficient in all sorts of NLP tasks out of the box.
- These extremely large scale Transformers are referred to these days as **Large Language Models (LLMs)**

# Transformers

- Previously we discussed the transformer architecture.
- Now, we'll see how this transformer can be trained to produce the LLMs we are familiar with.

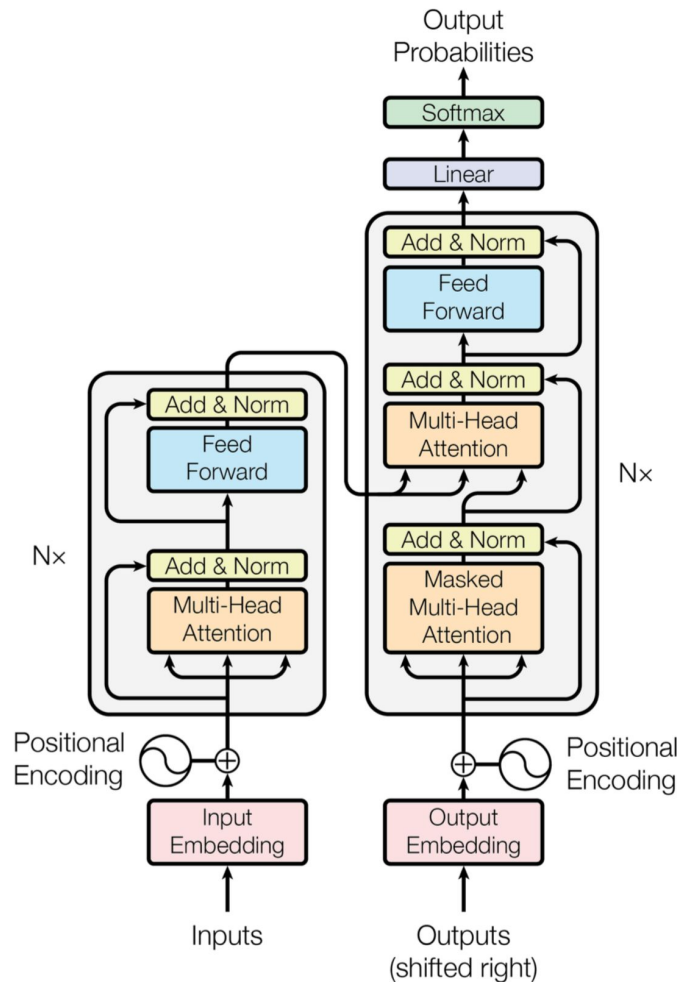


Figure 1: The Transformer - model architecture.

# Model type

- **Encoder-only**

- good for generating a representation of a document
- e.g. BERT, RoBERTa, Albert

- **Encoder-decoder**

- good for usages where output **heavily relies on input** (machine translation, text summarization, etc.)
- e.g. Flan-T5, BART

- **Decoder-only**

- good for text generation
- e.g. LLaMA series, Falcon, GPT series, LaMDA

# Model type

- **Encoder-only**

- good for generating a representation of a document
- e.g. BERT, RoBERTa, Albert

- **Encoder-decoder**

- good for usages where output **heavily relies on input** (machine

Since text generation applications are the most widely developed the past years, we'll focus on decoder-only architectures from now on

- **Decoder-only**

- good for text generation
- e.g. LLaMA series, Falcon, GPT series, LaMDA

# Contents

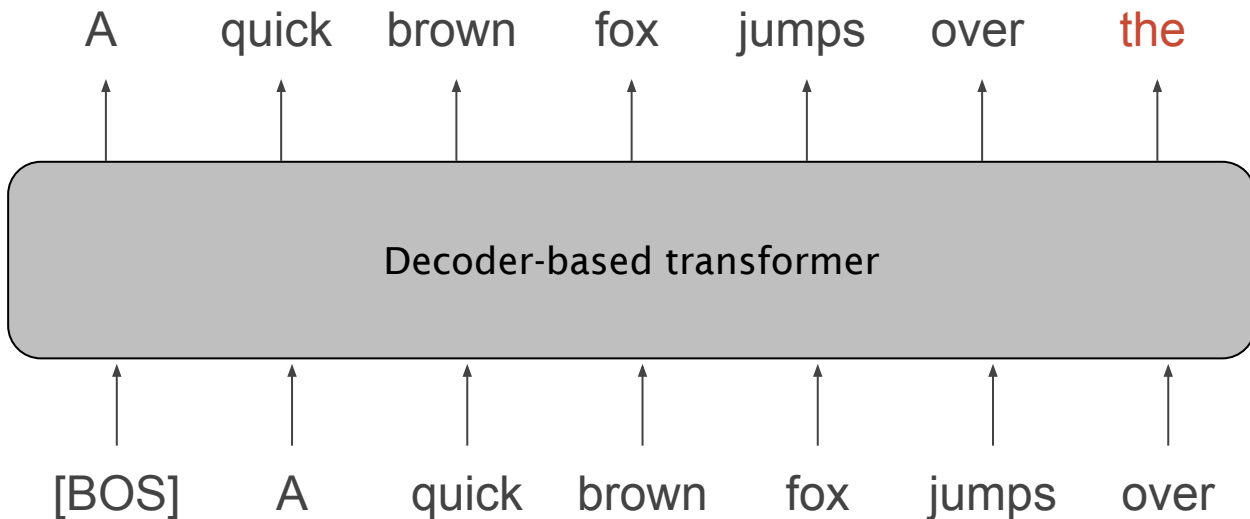
1. Training Language Models
  - I. Unsupervised pre-training
  - II. From LMs to assistants
    - A. Prompt Engineering
    - B. Supervised Fine-Tuning (SFT)
    - C. Reinforcement Learning from Human Feedback (RLHF)
2. Generating text
  - I. Decoding strategies

# Part 1. Training

## Unsupervised Pretraining

# Pre-training

- To pre-train a LLM a self-supervised learning technique is employed called **Language Modelling**
- The task is to simply **predict the next word in a sequence**





# Pre-training

- Note that training is inherently **unsupervised**, so we can use any data we can find!
- LMs are trained on lots of GB of data **scraped from the web**
  - e.g. wikipedia, quora, books.
- After training LMs can generate texts **autoregressively**
  - Start with an input prompt and ask the LM to predict the next token
  - Then append this predicted token to the end of the prompt and ask it to predict the next token
  - Repeat this process until an [EOS] token is generated

# So what does pretraining teach?

The Parthenon is a former temple located in \_\_\_\_\_, Greece.

Trivia

# So what does pretraining teach?

He sat in the front seat of \_\_\_\_ car.

Syntax

# So what does pretraining teach?

The man grabbed \_\_\_\_\_ arm in pain.

Coreference

# So what does pretraining teach?

The shelter was full of stray \_\_\_\_ and dogs.

Lexical semantics

# So what does pretraining teach?

I'm never going to watch another Tarantino movie again, the last movie I watched was \_\_\_\_\_.

Sentiment

# So what does pretraining teach?

... several key aspects of a language

That's why we call it **Language Modelling**!

# Part 1. Training

From LM to assistant



# Is pretraining enough?

- After pretraining our transformer is a proper Language Model, capable of understanding several key aspects of a language.
- What if we wanted our model to follow our instructions.

Prompt *Explain the moon landing to a 6 year old in a few sentences.*

# Is pretraining useful?

- In reality this is an illformed question from the beginning.
- Pretraining is extremely useful for producing a very good base model, which we can adapt to perform other tasks (one of which was the previous example)
- What we have after this intensive pretraining is a **foundational model**
  - GPT-3 is such a foundational model
  - It was adapted to follow user instructions to produce ChatGPT

# How to repurpose a LM as an assistant

- LMs are good at producing coherent output, but are **not inherently aligned with user intent**
- What we are more interested in is a **multitask assistant!**

Prompt: *Explain the moon landing to a 6 year old in a few sentences.*

Completion: GPT-3

Explain the theory of gravity to a 6 year old.

Explain the theory of relativity to a 6 year old in a few sentences.

Explain the big bang theory to a 6 year old.

Explain evolution to a 6 year old.


InstructGPT

People went to the moon, and they took pictures of what they saw, and sent them back to the earth so we could all see them.

# How to repurpose a LM as an assistant

- We know that LMs have learned the necessary information to be useful
- The only question is ...

... how do we go  
from **this** to **this**?



Prompt	<i>Explain the moon landing to a 6 year old in a few sentences.</i>
Completion	GPT-3
	Explain the theory of gravity to a 6 year old.
	Explain the theory of relativity to a 6 year old in a few sentences.
	Explain the big bang theory to a 6 year old.
	Explain evolution to a 6 year old.
	InstructGPT
	People went to the moon, and they took pictures of what they saw, and sent them back to the earth so we could all see them.

# How to repurpose a LM as an assistant

Three options:

## **A. Prompting**

- construct a prompt to extract relevant information from the LM

## **B. Supervised fine-tuning**

- continue the model's training in a supervised manner on pairs of (prompt, desired\_output)

## **C. Human feedback**

- use human feedback to evaluate the prompt's outputs and teach it to produce more desirable outputs

# Part 1. Training

From LM to assistant

A. Prompt Engineering

# Zero/Few shot capabilities of LLMs

- It has been shown that Large LMs exhibit latent **zero** and **few-shot capabilities**
- The latter is also referred to as **in-context learning**
- e.g. imagine writing the following prompt to a LLM

thanks -> merci

hello -> bonjour

goodbye ->

- The LLM would try to predict the most probable next tokens which in this case might be "au revoir"

# Chain-of-Thought

## Standard Prompting

### Model Input

Q: Roger has 5 tennis balls. He buys 2 more cans of tennis balls. Each can has 3 tennis balls. How many tennis balls does he have now?

A: The answer is 11.

Q: The cafeteria had 23 apples. If they used 20 to make lunch and bought 6 more, how many apples do they have?

### Model Output

A: The answer is 27. ❌

## Chain-of-Thought Prompting

### Model Input

Q: Roger has 5 tennis balls. He buys 2 more cans of tennis balls. Each can has 3 tennis balls. How many tennis balls does he have now?

A: Roger started with 5 balls. 2 cans of 3 tennis balls each is 6 tennis balls.  $5 + 6 = 11$ . The answer is 11.

Q: The cafeteria had 23 apples. If they used 20 to make lunch and bought 6 more, how many apples do they have?

### Model Output

A: The cafeteria had 23 apples originally. They used 20 to make lunch. So they had  $23 - 20 = 3$ . They bought 6 more apples, so they have  $3 + 6 = 9$ . The answer is 9. ✅



# Part 1. Training

From LM to assistant

B. Supervised Fine-Tuning (SFT)

# Supervised Fine-Tuning

- We have a very good model that doesn't do *exactly* what we want?
- What if we go to our bread and butter in Deep Learning: fine-tuning!
- Process:
  1. Collect examples of (prompt, desired output) over multiple tasks
    - can either pay people to write these
    - or look for such datasets in the web (e.g. questions/answers)
  2. Fine-tune the LM on these examples in a supervised manner
  3. Profit

# Issues with SFT

- Very promising approach on paper, simple and straightforward
- But...
  - no right answer on creative tasks (e.g. “write me a story about ...”) → hard to train for such tasks
  - not all mistakes are equally bad!

Let's say our (x, y) pair is:

("Where is the parthenon located?", "The parthenon is located in Athens.")

Model output 1: "The parthenon is located in Greece".

Model output 2: "The parthenon is located in Thessaloniki".

Both outputs would be penalized equally

- Mismatch between training objective and human preferences

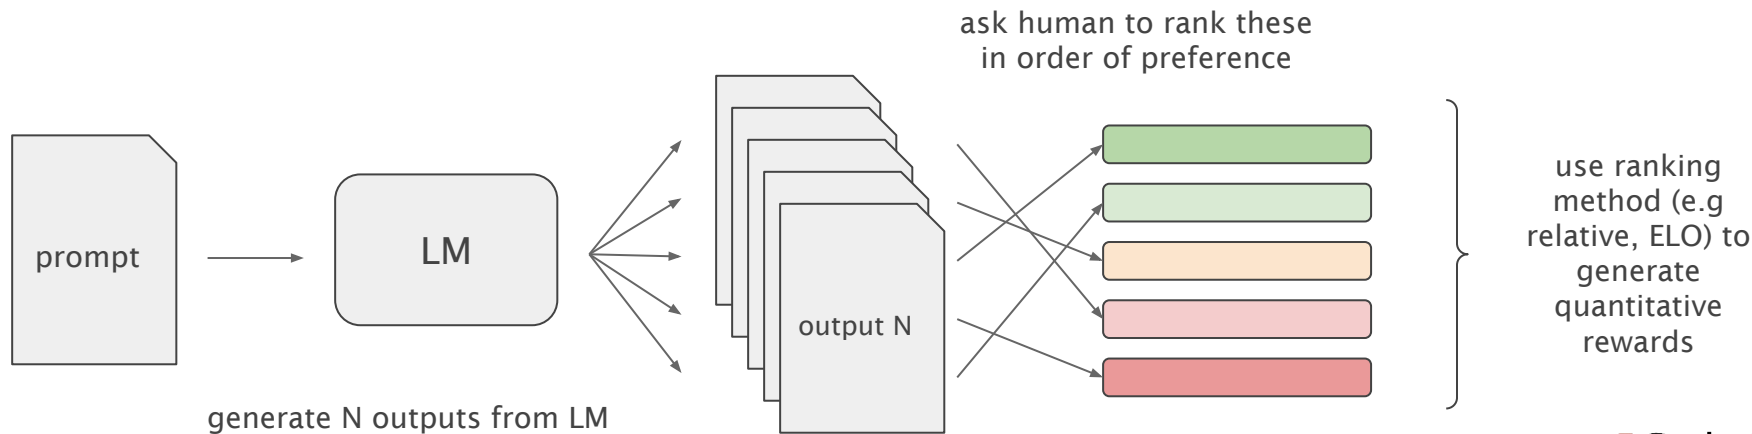
# Part 1. Training

From LM to assistant

C. Reinforcement Learning from Human Feedback  
(RLHF)


# Human Feedback

- Let's say we have some way of obtaining qualitative human feedback for a LM's outputs
- We could then optimize the LM on this feedback signal
- How do we get an reliable human feedback signal though?




# Reward Model

- Now that we have some (output, reward) pairs, we want to scale things up
- Train another LM (we'll call this the Reward Model) to predict the reward from a given output



# Reinforcement Learning

- The previously trained Reward Model (RM) is capable of predicting how a human would rate the output of a LM
- The only thing that remains is to fine-tune the LM on the human feedback
- We'll use an RL algorithm to train do this (usually PPO)



# RLHF recap

- RLHF Process:
  1. Collect RM training data
    - get some input prompts that interest us and ask our LM to produce several outputs for each
    - ask humans to rank these outputs in order of preference
    - obtain a quantitative reward from this ranking
  2. Train the RM
    - RM is another LM that accepts a text input and outputs a scalar value
    - train this in a supervised manner on the (output, reward) pairs of step 1
  3. Fine-tune the LM
    - generate outputs from input prompts
    - ask the RM to produce a reward from these outputs
    - use RL to train the LM to maximize reward



# Part 1. Training

From LM to assistant

Conclusion

# Which method to choose?

- We saw 3 possible ways of going from LM to assistant; each has their pros and cons
- Which one should we choose?
- These methods are **not** mutually exclusive!
- In fact modern chatbots (e.g. ChatGPT) were trained with all of these
  1. **SSL** for pretraining (i.e. language modelling)
  2. **SFT** afterwards to condition LM more towards instructions
  3. **RLHF** to make its outputs closer to human preferences
- Additionally, during inference we can use advanced **prompt engineering** techniques (e.g. chain of thought) to get even better results

# Foundational models


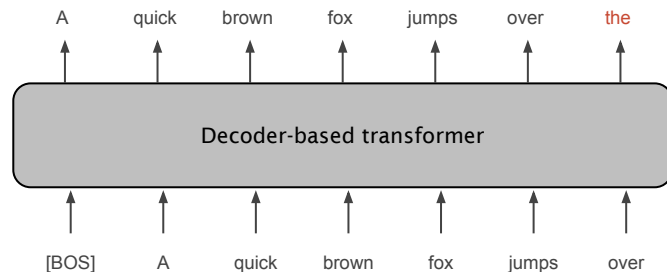
- We saw 3 possible ways of going from LM to assistant; each has their pros and cons
- Which one should we choose?
- These methods are **not** mutually exclusive!
- In fact modern chatbots (e.g. ChatGPT) were trained with all of these
  1. **SSL** for pretraining (i.e. language modelling)
  2. **SFT** afterwards to condition LM more towards instructions
  3. **RLHF** to make its outputs closer to human preferences
- Additionally, during inference we can use advanced **prompt engineering** techniques (e.g. chain of thought) to get even better results

# Part 1. Training

## Generating text

# How do LMs generate text?


- We've talked a lot about how LMs are trained but never actually discussed how they can generate text
- The figure to the right is a bit of an oversimplification
- In fact no LM outputs tokens; instead **they output a probability distribution over all tokens in the vocabulary**
- So how was the token “the” selected out of the mix?



# Decoding

The process of selecting the output tokens

- Does not change trainable parameters
- Can have noticeable effect in quality of outputs
- How the tokens will be selected based on their probabilities is referred to as the **decoding strategy**



# Decoding strategies

- **Greedy search:** select the token with the highest probability
- **Beam search:** keep the most likely  $N$  hypotheses; choose the one with the highest overall probability
- **Sampling:** randomly pick a word according to its probability
- **Top-k sampling:** similar to above, but only consider top-k tokens
- **Top-p (nucleus) sampling:** sample from the smallest set of words whose cumulative probability exceeds  $p$
- **Contrastive search**
- ... and more